

GEMM/GEMV in MLLM

A: q8_0(BSHD), B:q4_0x4(BSHD) -> C:f32(BSHD)

A: q8_0x4(BSHD), B:q4_0x4(BSHD) -> C:f32(BSHD)

for ARMv8.2+ devices

chenghua.wang.edu@gmail.com

The Quantization Data Types in llama.cpp


Q8_0	GH	8-bit round-to-nearest quantization (q). Each block has 32 weights. Weight formula: $w = q * \text{block_scale}$. Legacy quantization method (not used widely as of today).
Q8_1	GH	8-bit round-to-nearest quantization (q). Each block has 32 weights. Weight formula: $w = q * \text{block_scale} + \text{block_minimum}$. Legacy quantization method (not used widely as of today).
Q8_K	GH	8-bit quantization (q). Each block has 256 weights. Only used for quantizing intermediate results. All 2-6 bit dot products are implemented for this quantization type. Weight formula: $w = q * \text{block_scale}$.
I8	GH	8-bit fixed-width integer number.
Q6_K	GH	6-bit quantization (q). Super-blocks with 16 blocks, each block has 16 weights. Weight formula: $w = q * \text{block_scale}(8\text{-bit})$, resulting in 6.5625 bits-per-weight.
Q5_0	GH	5-bit round-to-nearest quantization (q). Each block has 32 weights. Weight formula: $w = q * \text{block_scale}$. Legacy quantization method (not used widely as of today).
Q5_1	GH	5-bit round-to-nearest quantization (q). Each block has 32 weights. Weight formula: $w = q * \text{block_scale} + \text{block_minimum}$. Legacy quantization method (not used widely as of today).
Q5_K	GH	5-bit quantization (q). Super-blocks with 8 blocks, each block has 32 weights. Weight formula: $w = q * \text{block_scale}(6\text{-bit}) + \text{block_min}(6\text{-bit})$, resulting in 5.5 bits-per-weight.
Q4_0	GH	4-bit round-to-nearest quantization (q). Each block has 32 weights. Weight formula: $w = q * \text{block_scale}$. Legacy quantization method (not used widely as of today).

```
typedef struct {  
    mllm_fp16_t d;           // delta  
    int8_t qs[QK8_0]; // quants QK8_0 = 32  
} block_q8_0;
```

```
// QK4_0 = 32  
typedef struct {  
    mllm_fp16_t d;           // delta  
    uint8_t qs[QK4_0 / 2]; // nibbles / quants  
} block_q4_0;
```

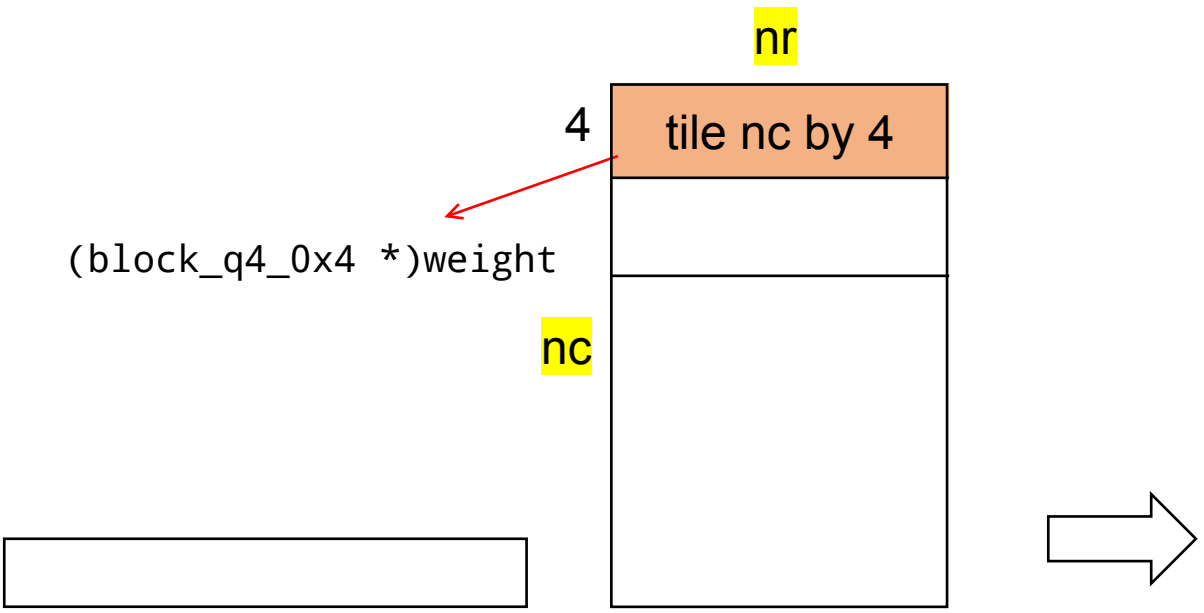
```
typedef struct {  
    mllm_fp16_t d[4];           // deltas for 4 q4_0 blocks  
    uint8_t qs[QK4_0 * 2]; // nibbles / quants for 4 q4_0 blocks  
} block_q4_0x4;
```

The GEMV impl in llama.cpp




When doing V Projection. The input has shape `[1, in_feature]`, and the weight's shape is `[out_feature, in_feature]`. Due to the quantization of weights is `q4_0x4`, we tile the column dimension of the weight matrix by 4.

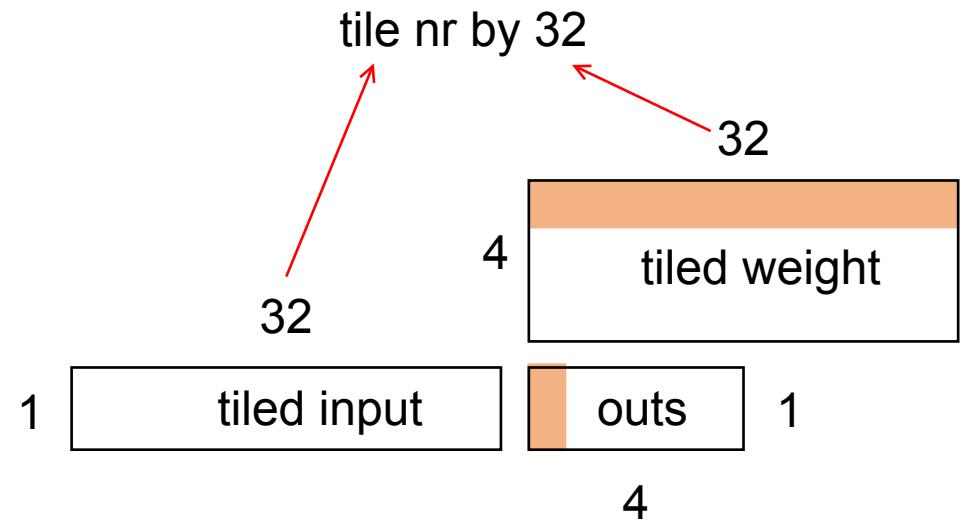
The GEMV impl in llama.cpp



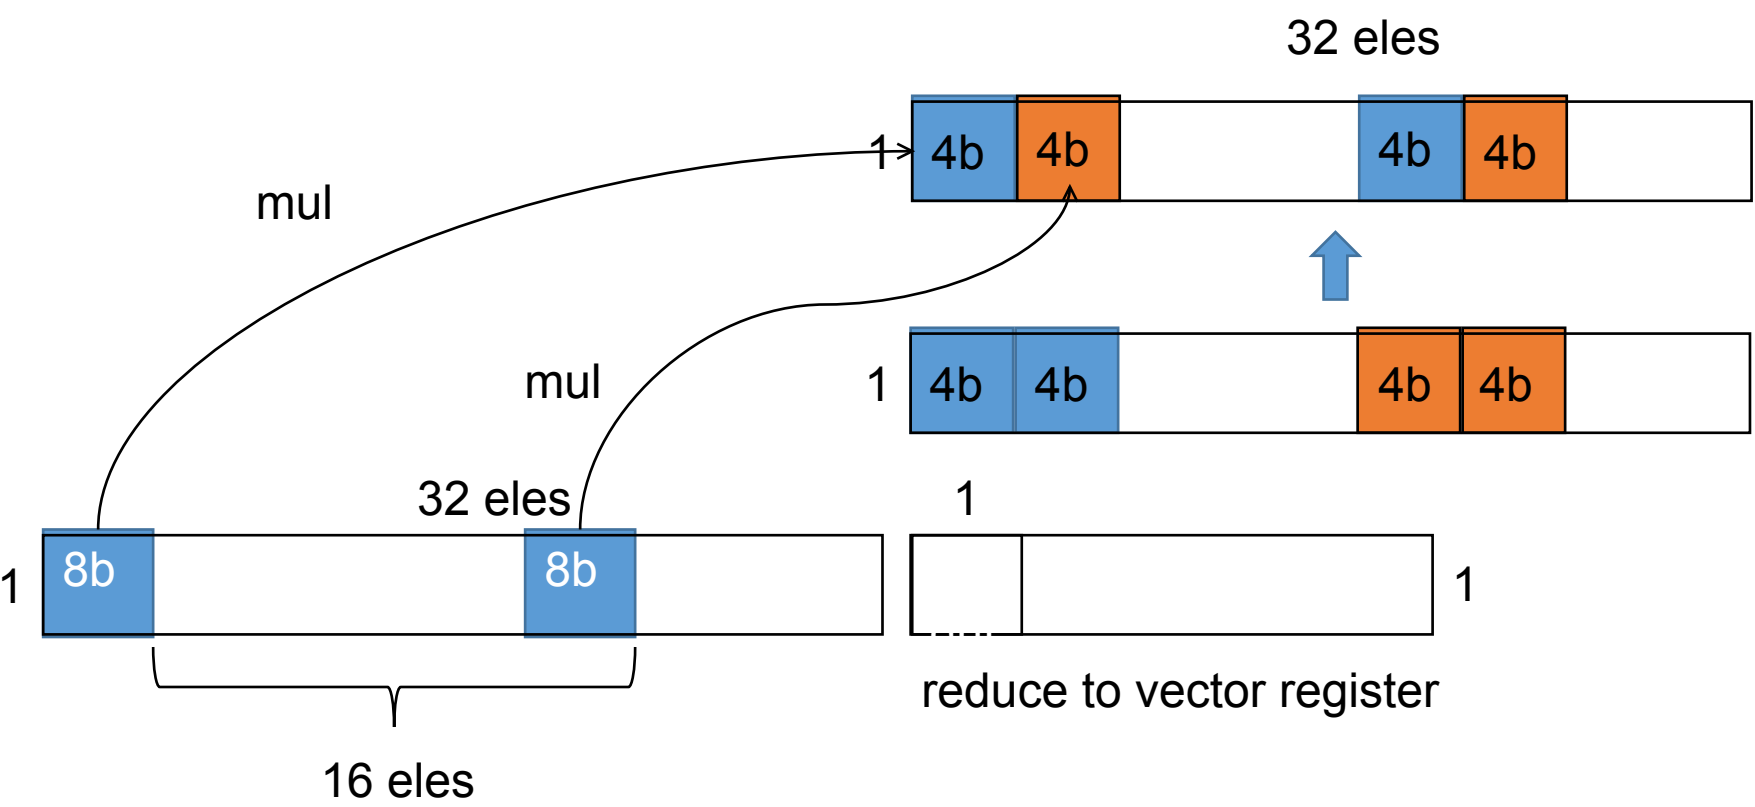
A tiled block([4, nr]) is an array of block_q4_0x4



```
for (int x = 0; x < nc / ncols_interleaved; x++)
```




The GEMV impl in llama.cpp





The GEMV ASM impl in llama.cpp

```
"movi v31.16b, #0x4\n"           // for sshl. to get high bits.
"movi v30.16b, #0xf0\n"          // for mask. to get low bits.
"add %[b_ptr], %[b_ptr], #0x8\n"  // to qs
"1:"                               // Column loop
"add x22, %[a_ptr], #0x2\n"        // to qs
"movi v29.16b, #0x0\n"            // acc is on register v29(16x8bits). Set to 0.
"mov x21, %[nb]\n"                 // move num of blocks to register x21
"2:"                               // Block loop
"ldr q28, [%x[b_ptr], #0x0]\n"     // load 128 bits from b matrix
"ldr q27, [x22, #0x0]\n"           // load 128 bits from a matrix
"movi v26.4s, #0x0\n"              // acc is on register v26(4x32bits). Set to 0.
"sub x20, x22, #0x2\n"             // to get scalar
"ldr q25, [x22, #0x10]\n"          // load 128 bits to q25. offsets is 16B
"ldr q24, [%x[b_ptr], #0x10]\n"    // load 128 bits to q24. offsets is 16B
"sub x21, x21, #0x1\n"             // nb = nb - 1
"add x22, x22, #0x22\n"            // a_ptr = aptr + 34B
"ldr q23, [%x[b_ptr], #0x20]\n"    // load 128 bits to q23. offset is 32
"ldr q22, [%x[b_ptr], #0x30]\n"    // load 128 bits to q22. offset is 48
"ld1r { v21.8h }, [x20]\n"         // scalar 4x16bit
"ldr q20, [%x[b_ptr], #-0x8]\n"    // scalar 1x16bit
"sshl v16.16b, v28.16b, v31.16b\n" // get high bits in q4_0x4
"and v28.16b, v28.16b, v30.16b\n" // get low bits in q4_0x4
```



The GEMV ASM impl in llama.cpp

```
"sshl v19.16b, v24.16b, v31.16b\n" // get high bits in q4_0x4
"and v24.16b, v24.16b, v30.16b\n" // get low bits in q4_0x4
"add %[b_ptr], %[b_ptr], #0x48\n" // b_ptr = b_ptr + 72
"sshl v18.16b, v23.16b, v31.16b\n" // get high bits in q4_0x4
"and v23.16b, v23.16b, v30.16b\n" // get low bits in q4_0x4
".inst 0x4f9be21a // sdot v26.4s, v16.16b, v27.4b[0]\n"
"sshl v17.16b, v22.16b, v31.16b\n" // get high bits in q4_0x4
"and v22.16b, v22.16b, v30.16b\n" // get low bits in q4_0x4
"fcvtl v21.4s, v21.4h\n" // cvt 8x16b to 4x32b, scalar of a matrix
"fcvtl v16.4s, v20.4h\n" // cvt 8x16b to 4x32b, scalar of b matrix. reuse v16 register
".inst 0x4f99e39a // sdot v26.4s, v28.16b, v25.4b[0]\n"
"fmul v16.4s, v16.4s, v21.4s\n" // v16 = v16 * v21, scalar a * scalar b
".inst 0x4fbbe27a // sdot v26.4s, v19.16b, v27.4b[1]\n" // v19(8 bits) + v27(32bit, 1B) to v26(32bit)
".inst 0x4fb9e31a // sdot v26.4s, v24.16b, v25.4b[1]\n"
".inst 0x4f9bea5a // sdot v26.4s, v18.16b, v27.4b[2]\n"
".inst 0x4f99eafa // sdot v26.4s, v23.16b, v25.4b[2]\n"
".inst 0x4fbbea3a // sdot v26.4s, v17.16b, v27.4b[3]\n"
".inst 0x4fb9eada // sdot v26.4s, v22.16b, v25.4b[3]\n"
"scvtf v26.4s, v26.4s, #0x4\n" // cvt int to float. the #0x4 is scale factor
"fmla v29.4s, v26.4s, v16.4s\n" // v29 = v26 * v16 + v29
"cbnz x21, 2b\n" // is x21 is not zero, jmp to label 2. num block loop.
"sub %[nc], %[nc], #0x4\n" // sub col by 4
"str q29, [%x[res_ptr], #0x0]\n" // store value to res_ptr
"add %[res_ptr], %[res_ptr], #0x10\n" // res_ptr move 16B. 4xf32.
"cbnz %[nc], 1b\n" // if nc is not zero, jump to label 1. num col loop.
```



Tiled B

The GEMM impl in llama.cpp

